

Understanding Attention: From Signals to Learning

Berkay Guler

Center for Pervasive Communications and Computing
University of California, Irvine

October 28, 2025

Outline

1. Motivation: The Problem We are Solving
2. Core Intuition: Adaptive Weighted Combinations
3. Scaling Attention
4. Multi-Head Attention: Parallel Processing
5. Self-Attention: When Everything Comes from One Source
6. Why Attention Works: A Computational View
7. Applications Across Domains
8. Limitations and Conclusions

Goal: Build understanding, not memorize formulas.

Part 1: Motivation

The Fundamental Problem

Classic Signal Processing: Given n signals x_1, x_2, \dots, x_n , compute:

$$y = \sum_{i=1}^n w_i x_i$$

where $x_i \in \mathbb{R}^d$ (signal vector), $w_i \in \mathbb{R}$ (scalar weight), $y \in \mathbb{R}^d$ (output).

Key observation: The weights w_i are usually *fixed* and predetermined.

But what if:

- ▶ The relevant signals depend on *context*?
- ▶ We want weights w_i that *adapt* based on what we're looking for?
- ▶ We want to learn which signals matter?

Reformulating the Problem

Instead of fixed weights, let's make them *data-dependent*:

$$w_i = f(\text{context}, x_i)$$

where:

- ▶ context $\in \mathbb{R}^d$: query vector (what we're looking for)
- ▶ $x_i \in \mathbb{R}^d$: the i -th signal vector
- ▶ $w_i \in \mathbb{R}$: scalar weight for signal i

Questions:

1. What is “context”? How do we represent it?
2. How should context and x_i interact to determine w_i ?
3. How do we ensure weights are meaningful (e.g., sum to 1)?
4. Can we compute this efficiently?

Part 2: Core Intuition

Three Representations: Query, Key, Value

Consider n elements with key-value pairs: $\{(k_1, v_1), (k_2, v_2), \dots, (k_n, v_n)\}$

- ▶ $i \in \{1, 2, \dots, n\}$: **Index** — element identifier
- ▶ $k_i \in \mathbb{R}^d$: **Key** — descriptor/fingerprint of element i
- ▶ $v_i \in \mathbb{R}^d$: **Value** — actual content of element i
- ▶ $q \in \mathbb{R}^d$: **Query** — what we're looking for

Goal: Weighted combination of values, where weights reflect key-query relevance.

Note: The value dimension can be different than that of queries and keys, i.e. d , but keys and queries must have the same dimension due to the dot product computation in the next page.

Intuition (information retrieval):

- ▶ Index i : Book number in library catalog
- ▶ Query q : “I’m searching for documents about machine learning”
- ▶ Key k_i : Document metadata/summary for book i
- ▶ Value v_i : Full document content of book i

Step 1: Compute Relevance Scores

How relevant is element i to our query? Measure via dot product:

$$s_i = q^T k_i \in \mathbb{R}$$

Why dot product?

- ▶ Correlation-based (from signal processing)
- ▶ s_i is large when q and k_i are aligned
- ▶ s_i is small when orthogonal/opposite
- ▶ Differentiable and efficient

Example scores:

$$s = [2.1, -0.5, 3.8, 0.2] \in \mathbb{R}^4$$

Element 3 is most relevant (highest score).

Step 2: Normalize via Softmax

Convert scores $s_i \in \mathbb{R}$ to weights $w_i \in \mathbb{R}$ via softmax:

$$w_i = \text{softmax}(s)_i = \frac{\exp(s_i)}{\sum_{j=1}^n \exp(s_j)}$$

Properties:

- ▶ $0 < w_i < 1$ for all i
- ▶ $\sum_i w_i = 1$ — probability distribution
- ▶ $w_i > w_j$ iff $s_i > s_j$ — preserves ordering
- ▶ Differentiable everywhere

Alternatives: Any function that normalizes to a probability distribution works:

- ▶ Sparsemax: produces sparse weights (some exactly zero)
- ▶ Sigmoid + normalize: $w_i = \frac{\sigma(s_i)}{\sum_j \sigma(s_j)}$
- ▶ Hardmax: $w_i = 1$ if $i = \arg \max_j s_j$, else $w_i = 0$ (not differentiable)

Softmax is standard due to smoothness and gradient properties.

Step 3: Compute Weighted Sum

Combine values using computed weights:

$$y = \sum_{i=1}^n w_i v_i \in \mathbb{R}^d$$

where $w_i \in \mathbb{R}$ (weight), $v_i \in \mathbb{R}^d$ (value), $y \in \mathbb{R}^d$ (output).

Complete process:

1. Score: $s_i = q^T k_i \in \mathbb{R}$
2. Weight: $w_i = \text{softmax}(s)_i \in \mathbb{R}$
3. Output: $y = \sum w_i v_i \in \mathbb{R}^d$

Result: An *adaptive superposition* of values, weighted by query-key relevance.

Example: Setup

Toy example: $d = 2$ (dimension), $n = 3$ (elements)

Query: $q = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \in \mathbb{R}^2$

Keys: $k_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad k_2 = \begin{bmatrix} 0.5 \\ 0.5 \end{bmatrix}, \quad k_3 = \begin{bmatrix} -1 \\ -1 \end{bmatrix}$

Values: $v_1 = \begin{bmatrix} 10 \\ 0 \end{bmatrix}, \quad v_2 = \begin{bmatrix} 0 \\ 10 \end{bmatrix}, \quad v_3 = \begin{bmatrix} 5 \\ 5 \end{bmatrix}$

Step 1: Scores $s_i = q^T k_i$

$$s_1 = 1, \quad s_2 = 1, \quad s_3 = -2$$

Score vector: $s = [1, 1, -2]^T \in \mathbb{R}^3$

Example: Softmax

Step 2: Softmax

Exponentials:

$$\exp(1) \approx 2.718, \quad \exp(-1) \approx 2.718, \quad \exp(-2) \approx 0.135$$

Partition: $Z = 2.718 + 2.718 + 0.135 = 5.571$

Weights:

$$w_1 = \frac{2.718}{5.571} \approx 0.488, \quad w_2 \approx 0.488, \quad w_3 \approx 0.024$$

Check: $0.488 + 0.488 + 0.024 = 1.000$

Weight vector: $w = [0.488, 0.488, 0.024]^T$

Example: Output

Step 3: Weighted sum

$$\text{output} = w_1 v_1 + w_2 v_2 + w_3 v_3$$

$$= 0.488 \begin{bmatrix} 10 \\ 0 \end{bmatrix} + 0.488 \begin{bmatrix} 0 \\ 10 \end{bmatrix} + 0.024 \begin{bmatrix} 5 \\ 5 \end{bmatrix}$$

$$= \begin{bmatrix} 4.88 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ 4.88 \end{bmatrix} + \begin{bmatrix} 0.12 \\ 0.12 \end{bmatrix} = \begin{bmatrix} 5.00 \\ 5.00 \end{bmatrix}$$

Interpretation: Output is dominated by v_1 and v_2 (weights 0.488 each), element 3 barely contributes.

Key Insight: Adaptivity

Same data, different query: $q' = \begin{bmatrix} -1 \\ -1 \end{bmatrix}$

Scores:

$$s'_1 = -1, \quad s'_2 = -1, \quad s'_3 = 2$$

After softmax:

$$w'_1 \approx 0.045, \quad w'_2 \approx 0.045, \quad w'_3 \approx 0.909$$

Now element 3 dominates!

Key point: Same data, different queries yield different attention patterns. This adaptivity makes attention powerful.

Attention: Compact Formula

Given:

- ▶ Query $q \in \mathbb{R}^d$
- ▶ Keys $K = [k_1, \dots, k_n] \in \mathbb{R}^{d \times n}$
- ▶ Values $V = [v_1, \dots, v_n] \in \mathbb{R}^{d \times n}$

Compact formula:

$$\text{Attention}(q, K, V) = V \cdot \text{softmax}(K^T q)$$

Expanded:

$$\text{output} = \sum_{i=1}^n \frac{\exp(q^T k_i)}{\sum_{j=1}^n \exp(q^T k_j)} v_i$$

Everything is differentiable and efficient via matrix operations.

Part 3: Scaling

The Scaling Issue

When d is large, dot products $q^T k_i$ have large magnitude.

Problem: Softmax becomes very peaked (one weight 1, others 0).

Softmax gradient: For $w_i = \frac{\exp(s_i)}{\sum_j \exp(s_j)}$, we have:

$$\frac{\partial w_i}{\partial s_i} = w_i(1 - w_i)$$

Derivation: $\frac{\partial}{\partial s_i} \left[\frac{\exp(s_i)}{Z} \right] = \frac{\exp(s_i) \cdot Z - \exp(s_i) \cdot \exp(s_i)}{Z^2} = w_i - w_i^2 = w_i(1 - w_i)$

When softmax is peaked ($w_i \approx 1$), gradient $\rightarrow 0$ (vanishing gradients).

Solution: Scale dot products by $1/\sqrt{d}$:

$$s_i = \frac{q^T k_i}{\sqrt{d}}$$

This keeps s_i in reasonable range regardless of d .

Why Scale by \sqrt{d} ?

Variance analysis:

Assume q, k_i have i.i.d. components with mean 0, variance 1.

The dot product is:

$$q^T k_i = \sum_{j=1}^d q_j k_{i,j}$$

Since components are independent:

$$\text{Var}(q^T k_i) = \sum_{j=1}^d \text{Var}(q_j k_{i,j}) = \sum_{j=1}^d \text{Var}(q_j) \text{Var}(k_{i,j}) = d$$

Therefore: $\text{Std}(q^T k_i) = \sqrt{d}$

After scaling: $\text{Var}\left(\frac{q^T k_i}{\sqrt{d}}\right) = \frac{d}{d} = 1$

Result: Scores have unit variance (standard deviation = 1) regardless of dimension d .

This keeps softmax outputs balanced rather than extremely peaked, which maintains healthy gradients for learning.

Extending to Multiple Queries (1/3)

So far: Single query vector $q \in \mathbb{R}^d$ produces single output.

$$y = \text{Attention}(q, K, V) = V \cdot \text{softmax}\left(\frac{K^T q}{\sqrt{d}}\right) \in \mathbb{R}^d$$

In practice: We often need to process m different queries **in parallel**.

Question: Can we compute attention for multiple queries efficiently?

Answer: Yes! Stack them into a matrix and use matrix operations.

Extending to Multiple Queries (2/3)

Stack queries: $Q = [q_1; q_2; \dots; q_m] \in \mathbb{R}^{m \times d}$ (each row is a query)

Matrix formulation:

$$Y = \text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right) V \in \mathbb{R}^{m \times d}$$

Output: $Y = [y_1; y_2; \dots; y_m] \in \mathbb{R}^{m \times d}$ (each row y_i is output for query q_i)

Dimensions:

- ▶ $QK^T \in \mathbb{R}^{m \times n}$: all query-key scores
- ▶ Softmax applied row-wise: each query gets its own distribution
- ▶ $Y \in \mathbb{R}^{m \times d}$: one output row per query row

Extending to Multiple Queries (3/3)

CRITICAL: Queries are processed independently!

Each query q_i gets its own output y_i with NO interaction between queries.

Analogy:

- ▶ **Correct:** Ask m questions → get m separate answers
- ▶ **Wrong:** Ask m questions → get 1 combined answer

Purpose: Computational efficiency only (parallel GPU processing).

Each query follows the same three steps: score → normalize → weighted sum.

Scaled Attention Formula

Scaled dot-product attention (standard):

$$Y = \text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d}} \right) V$$

where:

- ▶ $Y \in \mathbb{R}^{m \times d}$: m answers (one output per query)
- ▶ $Q \in \mathbb{R}^{m \times d}$: m queries
- ▶ $K \in \mathbb{R}^{n \times d}$: n keys
- ▶ $V \in \mathbb{R}^{n \times d}$: n values
- ▶ \sqrt{d} : scaling factor

Dimensions:

- ▶ $QK^T \in \mathbb{R}^{m \times n}$: query-key scores

Part 4: Multiple Heads

Motivation for Multiple Heads

Problem: One attention operation may not capture all relationships.

Library analogy: Suppose you are researching "neural networks"

Single attention head: One search strategy

- ▶ Search by title keywords only
- ▶ Might miss books categorized differently

Multiple attention heads: Different search strategies in parallel

- ▶ Head 1: Search by title keywords
- ▶ Head 2: Search by author expertise
- ▶ Head 3: Search by publication date (recent work)
- ▶ Head 4: Search by citation patterns (influential work)

Each head uses different Keys (different ways to describe/index the same books).

Result: Combine insights from all search strategies for comprehensive results.

What is a "Head"?

Definition: A "head" is one complete attention operation with its own learned projections.

Single head attention:

- ▶ Learn: $W_Q \in \mathbb{R}^{d \times d_k}$, $W_K \in \mathbb{R}^{d \times d_k}$, $W_V \in \mathbb{R}^{d \times d_v}$
- ▶ Project: $Q = X_Q W_Q \in \mathbb{R}^{m \times d_k}$, $K = X_K W_K \in \mathbb{R}^{n \times d_k}$, $V = X_V W_V \in \mathbb{R}^{n \times d_v}$
- ▶ Compute: $\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \in \mathbb{R}^{m \times d_v}$

Multi-head attention: Run h attention operations in parallel, each head $i \in \{1, \dots, h\}$:

- ▶ Head i projects with: $W_Q^{(i)} \in \mathbb{R}^{d \times d_k}$, $W_K^{(i)} \in \mathbb{R}^{d \times d_k}$, $W_V^{(i)} \in \mathbb{R}^{d \times d_v}$
- ▶ Head i output: $Y_i = \text{Attention}(X_Q W_Q^{(i)}, X_K W_K^{(i)}, X_V W_V^{(i)}) \in \mathbb{R}^{m \times d_v}$

Key idea: Different $W^{(i)}$ matrices \rightarrow different "views" \rightarrow capture different relationships.

Multi-Head: Concatenation and Output

After computing all heads: We have h outputs, each of dimension $m \times d_v$

$$Y_1 \in \mathbb{R}^{m \times d_v}, \quad Y_2 \in \mathbb{R}^{m \times d_v}, \quad \dots, \quad Y_h \in \mathbb{R}^{m \times d_v}$$

Step 1: Concatenate all head outputs horizontally:

$$\text{Concat} = [Y_1 \mid Y_2 \mid \dots \mid Y_h] \in \mathbb{R}^{m \times (h \cdot d_v)}$$

where \mid denotes horizontal concatenation.

Step 2: Apply output projection to map back to dimension d :

$$\text{MultiHead}(X_Q, X_K, X_V) = \text{Concat} \cdot W_O \in \mathbb{R}^{m \times d}$$

where $W_O \in \mathbb{R}^{(h \cdot d_v) \times d}$ is a learned output projection matrix.

Note: The choices of d_k and d_v are design decisions. Common choice: $d_k = d_v = d/h$ for computational balance. Note that $h \cdot d_v = d$ keeps the concatenated dimension equal to the input dimension, but this is not required due to the output projection W_O .

Multi-Head Attention: Summary

Complete formula:

$$\text{MultiHead}(X_Q, X_K, X_V) = \text{Concat}(Y_1, \dots, Y_h)W_O$$

where:

$$Y_i = \text{Attention}(X_Q W_Q^{(i)}, X_K W_K^{(i)}, X_V W_V^{(i)})$$

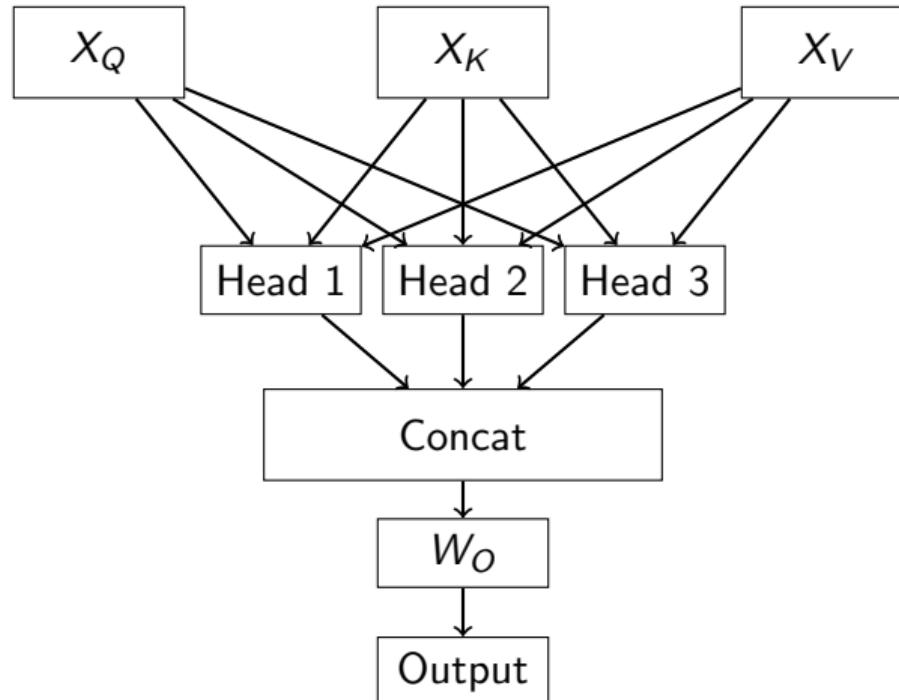
Learned parameters for h heads:

- ▶ h sets of projection matrices: $\{W_Q^{(i)} \in \mathbb{R}^{d \times d_k}, W_K^{(i)} \in \mathbb{R}^{d \times d_k}, W_V^{(i)} \in \mathbb{R}^{d \times d_v}\}_{i=1}^h$
- ▶ One output projection: $W_O \in \mathbb{R}^{(h \cdot d_v) \times d}$

Key properties:

- ▶ Each head operates independently (parallel computation)
- ▶ Each head output dimension is d_v (commonly $d_v = d/h$)
- ▶ Query/key dimensions are d_k (commonly $d_k = d/h$)
- ▶ More expressive: multiple different "views" of the relationships

Multi-Head Flow



Part 5: Self-Attention

Self-Attention: Everything from a Single Input

So far: We had separate inputs $X_Q \in \mathbb{R}^{m \times d}$, $X_K \in \mathbb{R}^{n \times d}$, $X_V \in \mathbb{R}^{n \times d}$.

Self-Attention: All three come from the same input $X \in \mathbb{R}^{n \times d}$:

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

where $W_Q, W_K \in \mathbb{R}^{d \times d_k}$ and $W_V \in \mathbb{R}^{d \times d_v}$ are learned projection matrices.

Note: Since $X_Q = X_K = X_V = X$, we have $m = n$ (number of queries equals number of keys/values).

Interpretation:

- ▶ Input X has n rows (elements), each of dimension d
- ▶ Each row serves as a query (via projection W_Q)
- ▶ All n rows serve as keys (via projection W_K) and values (via projection W_V)
- ▶ Output: Each row i attends to all rows $j = 1, \dots, n$
- ▶ Result: Each row gets an updated representation based on all other rows

Key insight: All n rows communicate with each other simultaneously.

Self-Attention Example (1/2)

Setup: $n = 3$ rows (elements), dimension $d = 2$

Input: $X = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{bmatrix} \in \mathbb{R}^{3 \times 2}$

Assume $W_Q = W_K = W_V = I_2$ (identity), so $Q = K = V = X$

Compute scores: $QK^T = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 1 & 2 \end{bmatrix}$

Scaled: $\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 1 & 2 \end{bmatrix} \approx \begin{bmatrix} 0.707 & 0 & 0.707 \\ 0 & 0.707 & 0.707 \\ 0.707 & 0.707 & 1.414 \end{bmatrix}$

Observation:

- ▶ Entry (i, j) captures the similarity between query i and key j
- ▶ Row 3 (which is $[1, 1]$) has highest score with itself (1.414)
- ▶ Row 3 also attends strongly to rows 1 and 2 (scores 0.707 each)
- ▶ Rows 1 and 2 are orthogonal (score 0)

Self-Attention Example (2/2)

Apply softmax row-wise:

Row 1: $\exp([0.707, 0, 0.707]) \rightarrow \text{normalize} \rightarrow w_1 \approx [0.401, 0.198, 0.401]$

Row 2: $\exp([0, 0.707, 0.707]) \rightarrow \text{normalize} \rightarrow w_2 \approx [0.198, 0.401, 0.401]$

Row 3: $\exp([0.707, 0.707, 1.414]) \rightarrow \text{normalize} \rightarrow w_3 \approx [0.248, 0.248, 0.504]$

Attention weights: $W_{\text{attn}} \approx \begin{bmatrix} 0.401 & 0.198 & 0.401 \\ 0.198 & 0.401 & 0.401 \\ 0.248 & 0.248 & 0.504 \end{bmatrix}$

Output: $Y = W_{\text{attn}} V = \begin{bmatrix} 0.401 & 0.198 & 0.401 \\ 0.198 & 0.401 & 0.401 \\ 0.248 & 0.248 & 0.504 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{bmatrix} \approx \begin{bmatrix} 0.802 & 0.599 \\ 0.599 & 0.802 \\ 0.752 & 0.752 \end{bmatrix}$

Note: W_{attn} is not symmetrical in general. Here, it is symmetrical since we have used the same projection I_2 for queries, keys, and values. In practice, we learn different projections for each.

Interpretation: Information is being mixed across all rows.

Why Do We Need Keys? Why Not Just Query-Value?

Valid alternative: Compute $s_i = q^T v_i$, then $w_i = \text{softmax}(s)_i$, then $y = \sum w_i v_i$.
So why use separate Keys?

Reason 1: Decoupling matching from content

- ▶ Keys: optimized for computing relevance/similarity
- ▶ Values: optimized for information content to propagate
- ▶ Same input can be "easy to match" (via W_K) yet "rich in content" (via W_V)

Reason 2: Dimensional flexibility

- ▶ Can use $d_k \neq d_v$ (e.g., smaller d_k for efficiency in score computation)
- ▶ Scores are $O(n^2 d_k)$, output is $O(n^2 d_v)$ — can optimize each separately

Reason 3: Empirical performance

- ▶ Two learned transformations (W_K, W_V) more expressive than one
- ▶ Model learns: "match on these features, retrieve different features"
- ▶ Significant performance gains in practice

Causal (Masked) Attention

Problem: In causal generation (auto-regressive models), position t shouldn't see future ($> t$).

Solution: Mask future positions before softmax:

$$\tilde{s}_{i,j} = \begin{cases} \frac{q_i^T k_j}{\sqrt{d}} & \text{if } j \leq i \\ -\infty & \text{if } j > i \end{cases}$$

After softmax: $\exp(-\infty) = 0$, so future positions get zero weight.

Result: Position i only attends to $1, \dots, i$ (causal).

Causal Masking Example

Unmasked scores:

$$\begin{bmatrix} 0.5 & 1.2 & 0.8 \\ 0.3 & 0.9 & 1.1 \\ 0.2 & 0.7 & 0.4 \end{bmatrix}$$

Apply causal mask: Row i (time $t = i$) can only attend to rows $j \leq i$

After causal mask:

$$\begin{bmatrix} 0.5 & -\infty & -\infty \\ 0.3 & 0.9 & -\infty \\ 0.2 & 0.7 & 0.4 \end{bmatrix}$$

After softmax:

$$\begin{bmatrix} 1.0 & 0 & 0 \\ 0.24 & 0.76 & 0 \\ 0.15 & 0.49 & 0.36 \end{bmatrix}$$

- ▶ Row 1 ($t=1$): attends only to itself
- ▶ Row 2 ($t=2$): attends to rows 1-2
- ▶ Row 3 ($t=3$): attends to all rows 1-3

Part 6: Why Attention Works

Computational Advantage: Adaptivity

Fixed architectures with hard-coded inductive biases:

RNN inductive bias: Sequential processing, recent context matters most

- ▶ Always combines h_{t-1} (previous state) and x_t (current input) the same way
- ▶ Information from distant past must flow through all intermediate steps
- ▶ Assumes temporal locality: nearby elements are most relevant

CNN inductive bias: Spatial locality and translation invariance

- ▶ Always looks at fixed local neighborhood (e.g., 3×3 receptive field)
- ▶ Same filter applied everywhere (translation invariance)
- ▶ Global context requires stacking many layers

Attention: Data-dependent, learned relevance

- ▶ Each element looks at *whatever is relevant*, not just neighbors
- ▶ Relevant context learned from data, not assumed by architecture
- ▶ Direct connections between any pair of elements

Trade-off: More flexible, but requires more data to learn patterns that RNN/CNN get "for free."

Expressiveness

1. **Aggregation (weighted sum):** This is what attention does directly
 - ▶ Attention weights w_i determine contribution of each element
2. **Filtering (suppress irrelevant):** Learn to assign near-zero weights
 - ▶ If $q^T k_i$ is very negative, then $w_i \approx 0$ after softmax
3. **Routing (select subset):** Learn sharp attention on a few elements
 - ▶ Make one or a few w_i large (e.g., $w_1 = 0.69$, $w_2 = 0.29$, rest ≈ 0)
 - ▶ Since $\sum w_i = 1$, this concentrates attention on selected elements
4. **Sorting (attend in order):** Most complex—requires multiple layers
 - ▶ Example: Input sequence [5, 2, 8, 1], want to process in sorted order
 - ▶ Layer 1: Learn to encode "which element is smallest/largest"
 - ▶ Layer 2: Position 1 learns to attend most to element with value 1
 - ▶ Layer 3: Position 2 learns to attend most to element with value 2, etc.
 - ▶ Result: Each output position aggregates information in sorted order
 - ▶ Not physically reordering, but retrieving information as if sorted

Key insight: By learning appropriate Q, K, V projections, attention implements diverse computational patterns.

Optimization

Gradient flow: Fully differentiable through:

1. Softmax to scores
2. Dot products to Q , K , V
3. Learned projections W_Q , W_K , W_V

Scaling benefit: $1/\sqrt{d}$ keeps gradients stable.

Result: End-to-end training via backprop works effectively.

Efficiency: Parallelization

Complexity: For sequence length n :

- ▶ All pairwise scores: $O(n^2d)$
- ▶ Softmax and weighting: $O(n^2d)$
- ▶ Total: $O(n^2d)$ per layer

Parallelization:

- ▶ No recurrence: h_t doesn't depend on h_{t-1}
- ▶ Process entire sequence simultaneously
- ▶ Natural fit for GPU/TPU

Trade-off: Higher memory, but massive parallelism.

Part 7: Applications

Application 1: Language Modeling

Task: Predict next token given context.

Approach:

1. Embed tokens
2. Stack of multi-head self-attention with causal mask
3. Each layer: position attends to context
4. Output: logits over vocabulary

Why attention helps:

- ▶ Long-range dependencies (100+ tokens back)
- ▶ Dynamic context (learns what's relevant)
- ▶ Parallelization (train on thousands of tokens)

Examples: GPT, LLaMA, Claude, ...

Application 2: Vision Transformers

Apply attention to images:

1. Divide image into patches (e.g., 16×16)
2. Embed each patch (i.e. project to some other space)
3. Self-attention: each patch attends to all patches
4. Use output for classification/other tasks

Key insight: Patches are like tokens.

Advantage over CNNs: No hard-coded locality. Model learns spatial relationships.

Application 3: Graph Attention Networks (1/2)

Graph-structured data:

1. Each node has features
2. Each node attends to its neighbors
3. Compute attention weights for neighbors
4. Update: weighted combination

Setup: Graph with N nodes, node i has features $h_i \in \mathbb{R}^d$, neighborhood $\mathcal{N}(i)$

Key idea: Each node aggregates information from its neighbors using attention weights based on feature similarity.

Useful for: social networks, molecular graphs, and knowledge graphs.

Application 3: Graph Attention Networks (2/2)

Attention mechanism with Q, K, V projections:

Project features: $q_i = W_Q h_i$, $k_j = W_K h_j$, $v_j = W_V h_j$

Compute scores: $s_{i,j} = \frac{q_i^T k_j}{\sqrt{d_k}}$ for each neighbor $j \in \mathcal{N}(i)$

Normalize: $w_{i,j} = \frac{\exp(s_{i,j})}{\sum_{k \in \mathcal{N}(i)} \exp(s_{i,k})}$

Update: $h'_i = \sum_{j \in \mathcal{N}(i)} w_{i,j} v_j$ **Key differences from standard Transformers:**

- ▶ Node i only attends to neighbors $\mathcal{N}(i)$, not all nodes (sparse attention)
- ▶ Attention is constrained by graph structure
- ▶ $W_Q, W_K, W_V \in \mathbb{R}^{d_k \times d}$ are learned projection matrices

Application 4: Cross-Attention

Query from one source, Keys/Values from another.

Recall the library analogy: Keys and values come from the same source, whereas the queries are independent!

Example: Generate text describing an image (Image-to-Text)

- ▶ Image encoder → image features
- ▶ Language decoder (with causal attention for autoregressive generation) generates text
- ▶ Decoder queries: “What in image is relevant?”
- ▶ Image features serve as Keys/Values

General pattern: Align two modalities via attention.

Other examples:

- ▶ Machine translation: source language is K/V, target language is Q
- ▶ Visual question answering (QA): question is Q, image regions are K/V

Application 5: Control and Robotics

Transformers for decision-making:

- ▶ Policy network takes history of observations/actions
- ▶ Self-attention: each timestep attends to relevant past → **causal attention!**
- ▶ Output: action to take

Why useful:

- ▶ Long-horizon reasoning (current action depends on distant past)
- ▶ Composable (combine heterogeneous information)
- ▶ Scalable (same architecture for different tasks)

Example: Navigation, manipulation from demonstrations, etc.

Why Attention Appears Everywhere

Attention solves a fundamental problem:

How do I adaptively combine multiple information sources based on relevance?

Wherever you have:

1. Multiple information sources
2. Need to select/aggregate based on context
3. Benefit from learning what's relevant

→ Attention likely helps.

This generality is rare. It's a primitive operation many tasks need.

Part 8: Limitations and Conclusions

Computational Limitations

Main constraint: Quadratic complexity $O(n^2)$ in sequence length.

Consequences:

- ▶ Long documents ($n > 10,000$): memory prohibitive
- ▶ Real-time processing: may not afford full attention
- ▶ Very large graphs or dense graphs: all-to-all attention intractable

Active research:

- ▶ Sparse attention (only nearby positions)
- ▶ Linear attention (kernel approximations)
- ▶ Hierarchical attention (clusters)3

Each trades expressiveness for efficiency.

Interpretability Challenge

Question: Can we understand what attention is doing?

Positive: Attention weights are interpretable (tell you what was weighted).

Challenge:

- ▶ Many layers (once you stack multiple attention blocks sequentially as in transformers) and heads → complex interactions
- ▶ Weights might not directly explain behavior
- ▶ Multiple heads do different things
- ▶ Depth (increases as you stack multiple attention blocks): information mixes in complex ways

Status: More interpretable than other deep models, but full understanding remains open.

Data Requirements

Observation: Attention models need more data than constrained architectures.

Why:

- ▶ No built-in inductive bias (unlike CNNs)
- ▶ Must learn from data what relationships matter
- ▶ More parameters to train

Trade-off:

- ▶ High data regime: Transformers (the most popular attention based architecture) excel
- ▶ Low data regime: Constrained models are often better

Practical implication: For small datasets, consider simpler models or add inductive biases.

Key Takeaways

1. Core Mechanism:

- ▶ Attention = adaptive weighted combination
- ▶ Query-Key-Value separation enables flexibility
- ▶ Softmax normalization ensures interpretable weights

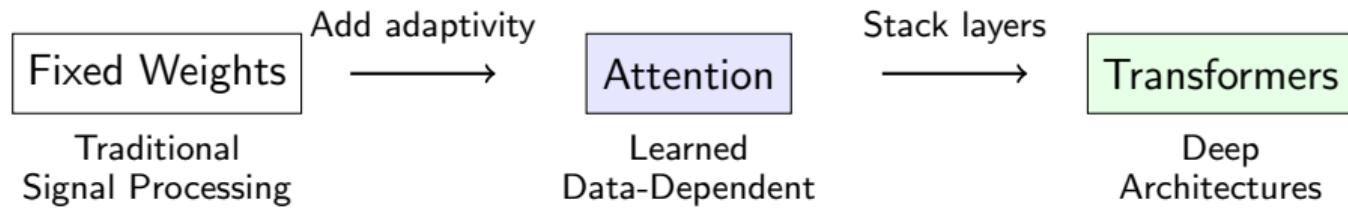
2. Why It Works:

- ▶ Adaptivity: learns what to attend to
- ▶ Expressiveness: can implement many operations
- ▶ Parallelizable: efficient on modern hardware

3. Universality:

- ▶ Not task-specific; a general primitive
- ▶ Appears in NLP, vision, graphs, control, ...
- ▶ Same mechanism, different applications

From Signal Processing to Learning



Journey: Fixed filtering → Adaptive attention → Deep learning architectures

Mathematical Summary

The Attention Operation:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d}} \right) V$$

Components:

- ▶ QK^T : Compute all query-key similarities
- ▶ $1/\sqrt{d}$: Scale for stability
- ▶ softmax: Normalize to probability distribution
- ▶ Multiply by V : Weighted combination of values

Multi-Head: Parallel attention with different learned projections, then concatenate and project.

Self-Attention: Q, K, V all from same input via learned projections.

Practical Advice

When to use attention:

- ▶ Multiple information sources to combine
- ▶ Context-dependent relevance
- ▶ Sufficient data and compute available
- ▶ Need long-range dependencies

When to be cautious:

- ▶ Very limited data (consider simpler models)
- ▶ Extremely long sequences (need approximations)
- ▶ Strong domain knowledge available (consider inductive biases)
- ▶ Interpretability critical (attention helps but isn't sufficient)

Implementation: Modern frameworks (PyTorch, JAX) have efficient implementations.
Start with standard architectures before customizing.

Resources for Further Study

Foundational papers:

- ▶ Vaswani et al. (2017): “Attention Is All You Need”
- ▶ Bahdanau et al. (2014): Original attention for NMT (Neural Machine Translation using attention in RNN-era)
- ▶ Dao et al. (2022): Flash Attention (efficient implementation)

Tutorials and courses:

- ▶ Stanford CS224N, MIT 6.S191
- ▶ 3Blue1Brown: Attention in transformers, step-by-step (Youtube Video)

Code implementations:

- ▶ PyTorch: `torch.nn.MultiheadAttention`
- ▶ Hugging Face Transformers library
- ▶ Annotated Transformer (Harvard NLP)

Conclusion

We've covered:

- ▶ Attention as adaptive weighted combination
- ▶ Mathematical formulation from first principles
- ▶ Scaling, multi-head extensions
- ▶ Why it works: adaptivity, expressiveness, efficiency
- ▶ Applications across domains
- ▶ Limitations

Attention is a powerful primitive for combining information.

Understanding it deeply enables you to apply it effectively and extend it creatively.

Questions?